

What's new in ClickHouse 20.12 - 21.1

[#15111](#) Implement gRPC protocol in ClickHouse.

[Vitaly Baranov](#) - Yandex.

The implementation of gRPC protocol also supports compression, SSL, getting progress and logs, authentication, parallel queries through the same channel, cancellation of queries, sessions, external tables.

```
/etc/clickhouse-server/config.xml
```

```
...
```

```
    <!-- <grpc_port>9100</grpc_port> -->
```

```
    <grpc>
```

```
        <enable_ssl>false</enable_ssl>
```

```
    <!-- The following two files are used only if enable_ssl=1 -->
```

```
    <ssl_cert_file>/path/to/ssl_cert_file</ssl_cert_file>
```

```
    <ssl_key_file>/path/to/ssl_key_file</ssl_key_file>
```

```
...
```

example (python) :

`https://github.com/ClickHouse/ClickHouse/blob/master/utils/grpc-client/clickhouse-grpc-client.py`

`utils/grpc-client/clickhouse-grpc-client.py -q "SELECT sum(number) FROM numbers(10)"`

`cat a.txt | utils/grpc-client/clickhouse-grpc-client.py -q "INSERT INTO temp FORMAT TSV"`

[#17144](#) *.zst compression/decompression support for data import and export.
It enables using *.zst in file() function and Content-encoding: zstd in HTTP client.

[Abi Palagashvili](#)

[#8526](#) support for brotli (br) compression in file-related storages and table functions.

[alexey-milovidov](#) - Yandex

```
CREATE TABLE file (x UInt64) ENGINE = File(TSV, 'data1.tsv.br');
```

```
SELECT count(), max(x)
FROM file('data{1,2}.tsv.{gz,br}', TSV, 'x UInt64');
```

```
$ echo "SELECT 1" | brotli | curl -sS --data-binary @- \
  -H 'Content-Encoding: br' "${CLICKHOUSE_URL}"
```

```
$ echo "SELECT 1" | zstd -c | curl -sS --data-binary @- \
  -H 'Content-Encoding: zstd' "${CLICKHOUSE_URL}";
```

Supported Content-Encoding: gzip, deflate, **br**, xz, **zstd**

[#15806](#) new data type Map.

First version for Map only supports String type of key and value. ().

[hexiaoting](#)

```
SET allow_experimental_map_type = 1;
```

```
create table table_map (a Map(String, String)) engine = Memory;
```

```
insert into table_map values  
  ({'name':'zhangsan', 'gender':'male'}),  
  ({'name':'lisi', 'gender':'female'});
```

```
SELECT a['name'] FROM table_map;
```

```
arrayElement(a, 'name')  
zhangsan  
lisi
```


[#16338](#) Implement UNION DISTINCT
and treat the plain UNION clause as UNION DISTINCT by default.

Add a setting `union_default_mode` that allows to treat it as
UNION ALL or require explicit mode specification. (DISTINCT / ALL)

[flynn](#)

```
SET union_default_mode = 'DISTINCT';
```

```
SELECT 'a' UNION SELECT 'a';
```

'a'
a

```
SET union_default_mode = 'ALL';
```

```
SELECT 'a' UNION SELECT 'a';
```

'a'
a

'a'
a

[#11617](#) Parallel formatting for data export.

[Nikita Mikhaylov](#) Yandex

```
time clickhouse-client  
-q 'select number, number/11111, toString(number) from numbers(100000000) format TSV' > /dev/null  
  
real    0m3.594s  
user    0m10.867s  
sys     0m0.308s
```

```
time clickhouse-client --output_format_parallel_formatting=0  
-q 'select number, number/11111, toString(number) from numbers(100000000) format TSV' > /dev/null  
  
real    0m9.410s  
user    0m9.128s  
sys     0m0.276s
```

[#7649](#) Added functions for calculation of minHash and simHash of text n-grams and shingles.

They are intended for semi-duplicate search.

Simhash, Minhash, bitHammingDistance, tupleHammingDistance

[flynn](#)

```
SELECT ngramSimHash('what a cute cat.');
```

```
SELECT ngramSimHashCaseInsensitive('what a cute cat.');
```

```
SELECT ngramSimHashUTF8('what a cute cat.');
```

```
SELECT ngramSimHashCaseInsensitiveUTF8('what a cute cat.');
```



```
SELECT wordShingleSimHash('what a cute cat.');
```

```
SELECT wordShingleSimHashCaseInsensitive('what a cute cat.');
```

```
SELECT wordShingleSimHashUTF8('what a cute cat.');
```

```
SELECT wordShingleSimHashCaseInsensitiveUTF8('what a cute cat.');
```



```
SELECT ngramMinHash('what a cute cat.');
```

```
SELECT ngramMinHashCaseInsensitive('what a cute cat.');
```

```
SELECT ngramMinHashUTF8('what a cute cat.');
```

```
SELECT ngramMinHashCaseInsensitiveUTF8('what a cute cat.');
```



```
SELECT wordShingleMinHash('what a cute cat.');
```

```
SELECT wordShingleMinHashCaseInsensitive('what a cute cat.');
```

```
SELECT wordShingleMinHashUTF8('what a cute cat.');
```

```
SELECT wordShingleMinHashCaseInsensitiveUTF8('what a cute cat.');
```



```
SELECT tupleHammingDistance((1, 2), (3, 4));
```

```
SELECT bitHammingDistance(100, 100000);
```

```
select ngramSimHash('Universal Avenue') a, ngramSimHash('Universe Avenue') b, a/b diff
```

a	b	diff
1577599152	1580943536	0.997884564550318

```
select ngramSimHash('Purple flowers') a, ngramSimHash('Green grass') b, a/b diff
```

a	b	diff
1485423320	3599596288	0.4126638659318453

[#16883](#) Added mannWitneyUTest, studentTTest and welchTTest aggregate functions. Refactored rankCorr a bit.

[Nikita Mikhaylov](#) - Yandex


```
SELECT rankCorr(number, number - 33)
FROM numbers(10)
```

```
rankCorr(number, minus(number, 33))
1
```

```
SELECT rankCorr(number, 1 / number)
FROM numbers(100)
```

```
rankCorr(number, divide(1, number))
-1
```

[#13403](#) ALTER UPDATE/DELETE IN PARTITION
with partition pruning in ReplicatedMergeTree

[Vladimir Chebotarev](#) - Altinity.

```
CREATE TABLE T (p Int64, d Date, s String)
ENGINE=ReplicatedMergeTree('/clickhouse/T', '{replica}')
PARTITION BY toYYYYMM(d) ORDER BY p;
```

```
INSERT INTO T VALUES(1, '2020-01-01', 'x'), (2, '2020-10-01', 'y');
```

```
SELECT *, _part FROM T;
```

p	d	s	_part
1	2020-01-01	x	202001_0_0_0

p	d	s	_part
2	2020-10-01	y	202010_0_0_0

```
ALTER TABLE T UPDATE s = 'z' WHERE p = 1 and d = '2020-01-01';
```

```
SELECT *, _part FROM T;
```

p	d	s	_part
1	2020-01-01	z	202001_0_0_0_1

p	d	s	_part
2	2020-10-01	y	202010_0_0_0_1

```
ALTER TABLE T UPDATE s = 'z'  
    IN PARTITION tuple(toYYYYMM(toDate('2020-01-01')))  
    WHERE p = 1;
```

```
SELECT *, _part FROM T;
```

p	d	s	_part
1	2020-01-01	z	202001_0_0_0_1

p	d	s	_part
2	2020-10-01	y	202010_0_0_0_0

[#17642](#) DETACH TABLE/VIEW ... PERMANENTLY

[filimonov](#) - Altinity.

so that after restarting the table does not reappear back automatically on restart (only by explicit request).

The table can still be attached back using the short syntax ATTACH TABLE.

```
DETACH table T PERMANENTLY;
```

```
service clickhouse-server restart
```

```
select * from T
```

```
DB::Exception: Table dw.T doesn't exist.
```

```
attach table T;
```

```
select * from T;
```

p	d	s
1	2020-01-01	x

[#17846](#) Extended OPTIMIZE ... DEDUPLICATE
syntax to allow explicit list of columns to check for duplicates on.

[Vasily Nemkov](#) - Altinity.


```
CREATE TABLE T(country String, city String, d Date,  
                some_coll Int64, some_col2 Int64)  
Engine=MergeTree Order by (country, city);
```

```
INSERT INTO T VALUES('FR', 'Leon', '2020-01-01', 33, 42);  
INSERT INTO T VALUES('FR', 'Leon', '2020-01-01', 32, 41);  
INSERT INTO T VALUES('FR', 'Leon', '2020-01-01', 35, 41);
```

```
OPTIMIZE TABLE T DEDUPLICATE  
  BY * EXCEPT (some_coll, some_col2);
```

```
SELECT * FROM T
```

country	city	d	some_coll	some_col2
FR	Leon	2020-01-01	33	42

```
OPTIMIZE TABLE table DEDUPLICATE; -- the old one

OPTIMIZE TABLE table DEDUPLICATE BY *; -- excludes MATERIALIZED columns

OPTIMIZE TABLE table DEDUPLICATE BY * EXCEPT colX;

OPTIMIZE TABLE table DEDUPLICATE BY * EXCEPT (colX, colY);

OPTIMIZE TABLE table DEDUPLICATE BY col1,col2,col3;

OPTIMIZE TABLE table DEDUPLICATE BY COLUMNS('column-matched-by-regex');

OPTIMIZE TABLE table DEDUPLICATE BY COLUMNS('column-matched-by-regex') EXCEPT colX;

OPTIMIZE TABLE table DEDUPLICATE BY COLUMNS('column-matched-by-regex') EXCEPT (colX, colY);
```

[#15511](#) ALTER TABLE ... DROP PART 'part_name'

[nvartolomei](#) - Cloudflare

```
CREATE TABLE T (p Int64, d Date, s String)
ENGINE=ReplicatedMergeTree('/clickhouse/T', '{replica}')
PARTITION BY toYYYYMM(d) ORDER BY p;

INSERT INTO T VALUES (1, '2020-01-01', 'x');
INSERT INTO T VALUES (2, '2020-01-02', 'y');
```

```
SELECT *, _part FROM T;
```

p	d	s	_part
2	2020-01-02	y	202001_1_1_0

p	d	s	_part
1	2020-01-01	x	202001_0_0_0

```
ALTER TABLE T DROP PART '202001_0_0_0';
```

```
SELECT *, _part FROM T;
```

p	d	s	_part
2	2020-01-02	y	202001_1_1_0

[#16895](#) Remove empty parts after they were pruned by TTL, mutation, or collapsing merge algorithm.

[Anton Popov](#) - Yandex

[#15073](#) Add EmbeddedRocksDB table engine (can be used for dictionaries).

[sundyli](#) - ByteDance

```
CREATE TABLE T1 ( P Int64, S String )  
ENGINE = EmbeddedRocksDB  
PRIMARY KEY P;
```

```
INSERT INTO T1 VALUES (1, '');
```

```
INSERT INTO T1 VALUES (1, 'updated');
```

```
SELECT * FROM T1;
```

P	S
1	updated


```
INSERT INTO T1 select number, toString(number)
from numbers(10 000 000);
```

```
select * from T1 where P = 99999;
```

P	S
99999	99999

1 rows in set. Elapsed: 0.001 sec.

```
INSERT INTO T1 VALUES (99999, 'updated');
1 rows in set. Elapsed: 0.001 sec.
```

```
select * from T1 where P = 99999;
```

P	S
99999	updated

1 rows in set. Elapsed: 0.001 sec.

[#16123](#) Add setting `aggregate_functions_null_for_empty` for SQL standard compatibility. This option will rewrite all aggregate functions in a query, adding `-OrNull` suffix to them.

[flynn](#) - Institute of Computing Technology, Beijing

```
CREATE TABLE T (p Int64, d Date, s String)
ENGINE=MergeTree
PARTITION BY toYYYYMM(d) ORDER BY p;

INSERT INTO T VALUES (1, '2020-01-01', 'x'), (2, '2020-10-01', 'y');

SELECT max(p) FROM T WHERE s = '';
```

```
max(p)
0
```

```
SELECT max(p) FROM T WHERE s = ''
SETTINGS aggregate_functions_null_for_empty = 1;
```

```
maxOrNull(p)
null
```

[#16575](#) Subqueries in WITH section (CTE) can reference previous subqueries in WITH section by their name.

[Amos Bird](#)

```
WITH  
T1 as (select number as c1 from numbers(10000) where number%2),  
T2 as (select max(c1) sm from T1),  
T3 as (select * from T1 where c1 = (select sm from T2))  
select * from T3
```

c1
9999

[#16253](#) Now we can safely prune partitions with exact match.

Useful case: Suppose table is partitioned by $\text{intHash64}(x) \% 100$ and the query has condition on $\text{intHash64}(x) \% 100$ verbatim, not on x .

[Amos Bird](#)

```
CREATE TABLE T (p Int64, d Date, tenant_id Int64, s String)
ENGINE=MergeTree
PARTITION BY (toYYYYMM(d), tenant_id % 100)
ORDER BY p;
```

```
SELECT count() FROM T WHERE (tenant_id = 42) AND (d = '1970-01-11')
```

count()
97

1 rows in set. Elapsed: 0.003 sec. Processed 30.10 thousand rows

```
select count() from T where identity(tenant_id) = 42 and d = '1970-01-11';
```

count()
97

1 rows in set. Elapsed: 0.010 sec. Processed 3.01 million rows

```
CREATE TABLE T (p Int64, d Date, tenant_id Int64,  
    part_id materialized tenant_id % 100,  
    s String)  
ENGINE=MergeTree  
PARTITION BY (toYYYYMM(d), part_id)  
ORDER BY p;
```

```
select count() from T where tenant_id = 42 and part_id = 42%100 and d = '1970-01-11';
```

count()
97

1 rows in set. Elapsed: 0.014 sec. Processed 30.10 thousand rows

[#18637](#) SimpleAggregateFunction in SummingMergeTree.
Now it works like AggregateFunction.

[Amos Bird](#)

```
CREATE TABLE T
(
    country String,
    city String,
    some_sum Int64,
    some_max SimpleAggregateFunction(max, Int32),
    some_gr SimpleAggregateFunction(groupUniqArrayArray, Array(String))
)
ENGINE = SummingMergeTree
ORDER BY (country, city);

INSERT INTO T VALUES ('FR', 'Leon', 1, 1, ['j']);
INSERT INTO T VALUES ('FR', 'Leon', 1, 10, ['k']);

OPTIMIZE TABLE T FINAL;
```

country	city	some_sum	some_max	some_gr
FR	Leon	2	10	['k','j']