



From ElasticSearch To ClickHouse, Log Analysis Practice in bilibili

Jiaqi Shu (束家麒)
bilibili INFRA

Tao Ling (凌涛)
bilibili OLAP





Agenda

- ❖ Log Analysis System based on ClickHouse
 - Challenge & problems
 - Why ClickHouse?
 - How we build
- ❖ Optimization & Customization of ClickHouse for Log Analysis
 - Challenges for ClickHouse
 - Overall Design
 - Implicit Map





Log Analysis System based on ClickHouse





Challenge & problems

- PB level data, Based on Elasticsearch, 1k+. nodes
- Write performance bottleneck
- Heavy cost on index
- Operation and maintenance costs





Why ClickHouse?

Requirements

- Low cost
- High performance
- Support semi-structure



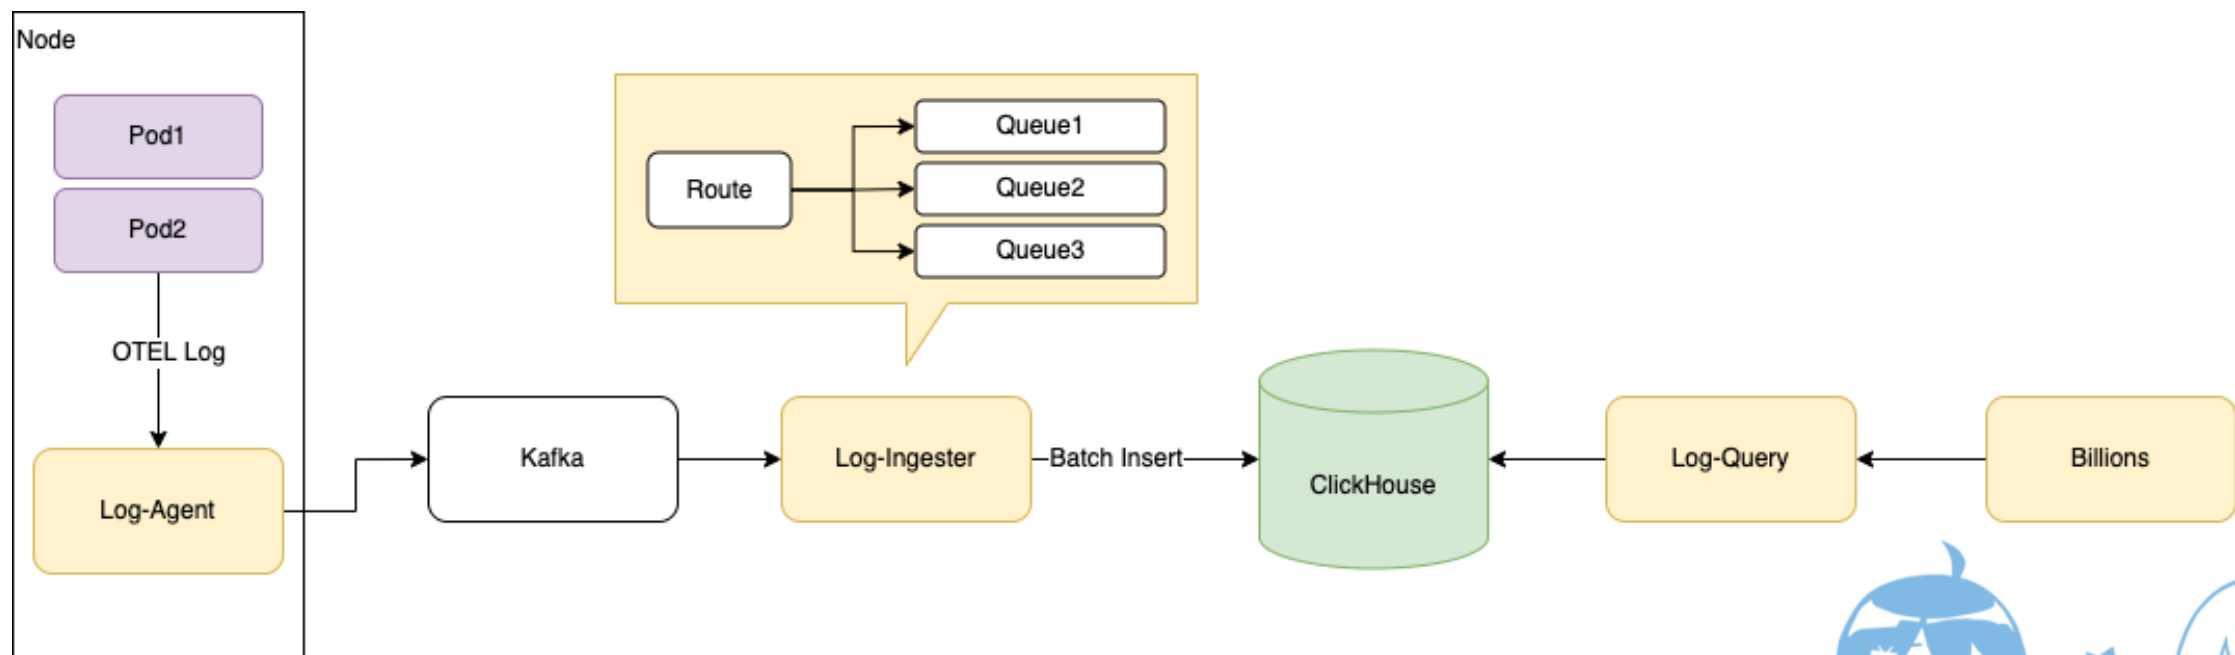
ClickHouse

- Cost saving over 60%
- Avg Query latency < 1s





How we build





How we build

```
Create Table <log_app_name>
{
  timestamp Datetime64,
  hostname  String CODC(ZSTD(1)),
  zone      String CODC(ZSTD(1)),
  log       String CODC(ZSTD(1)),
  ....
  string_map MapV2(String, Nullable (String))
                                   CODEC(ZSTD(1))
  number_map MapV2(String, Nullable (Float64))
                                   CODEC(ZSTD(1))
  bool_map   MapV2(String, Nullable (UInt8))
}
....
```

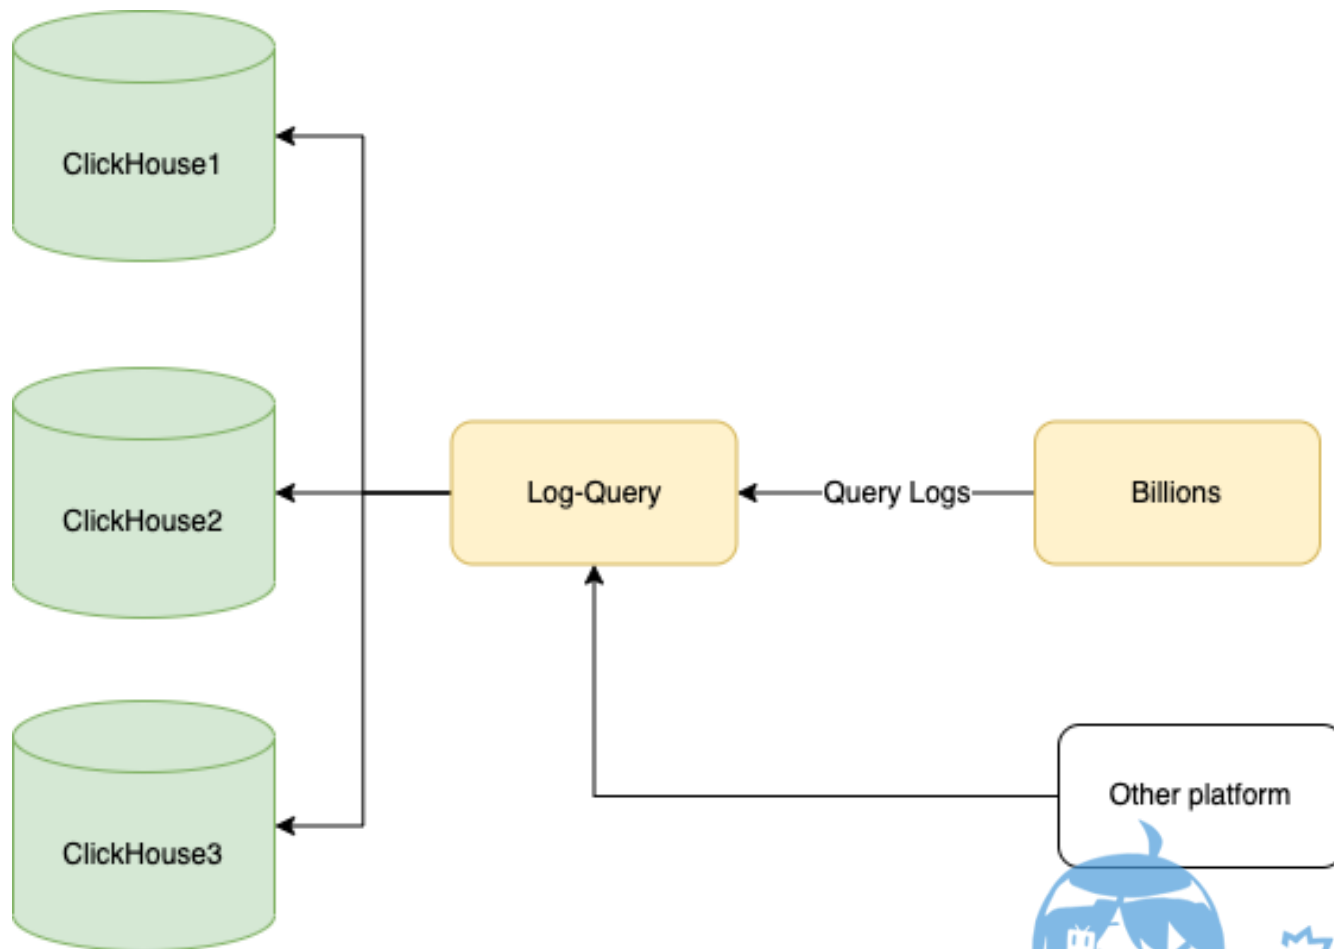




How we build

Log Query

- High level query
- Query routing
- Rate Limit, Access log...





How we build

```
{
  "app_id": "ops.billions.elastalert-worker",
  "query": "level = 'ERROR' AND cluster = 'shylf'",
  "from": "2022-03-02 06:00:00",
  "to": "2022-03-02 06:15:00"
}
```



```
SELECT *
FROM ops_billions_elastalert_worker
WHERE
  level = 'ERROR'
AND
  string_map['cluster'] = 'shylf'
AND
  timestamp >= '2022-03-02 06:00:00'
AND
  timestamp >= '2022-03-02 06:15:00'
ORDER BY
  timestamp desc
LIMIT 1000
```



timestamp	level	instance_name	log	string_map
2022-03-02 06:01:03	ERROR	shylf-opslog-02	fail to get connection	{"cluster": "shylf"}
2022-03-02 06:01:02	ERROR	shylf-opslog-01	io error	{"cluster": "shylf"}



```
[
  {
    "timestamp": "2022-03-02 06:01:03",
    "level": "ERROR",
    "log": "fail to get connection",
    "cluster": "shylf"
  },
  {
    "timestamp": "2022-03-02 06:01:02",
    "level": "ERROR",
    "log": "io error",
    "cluster": "shylf"
  }
]
```



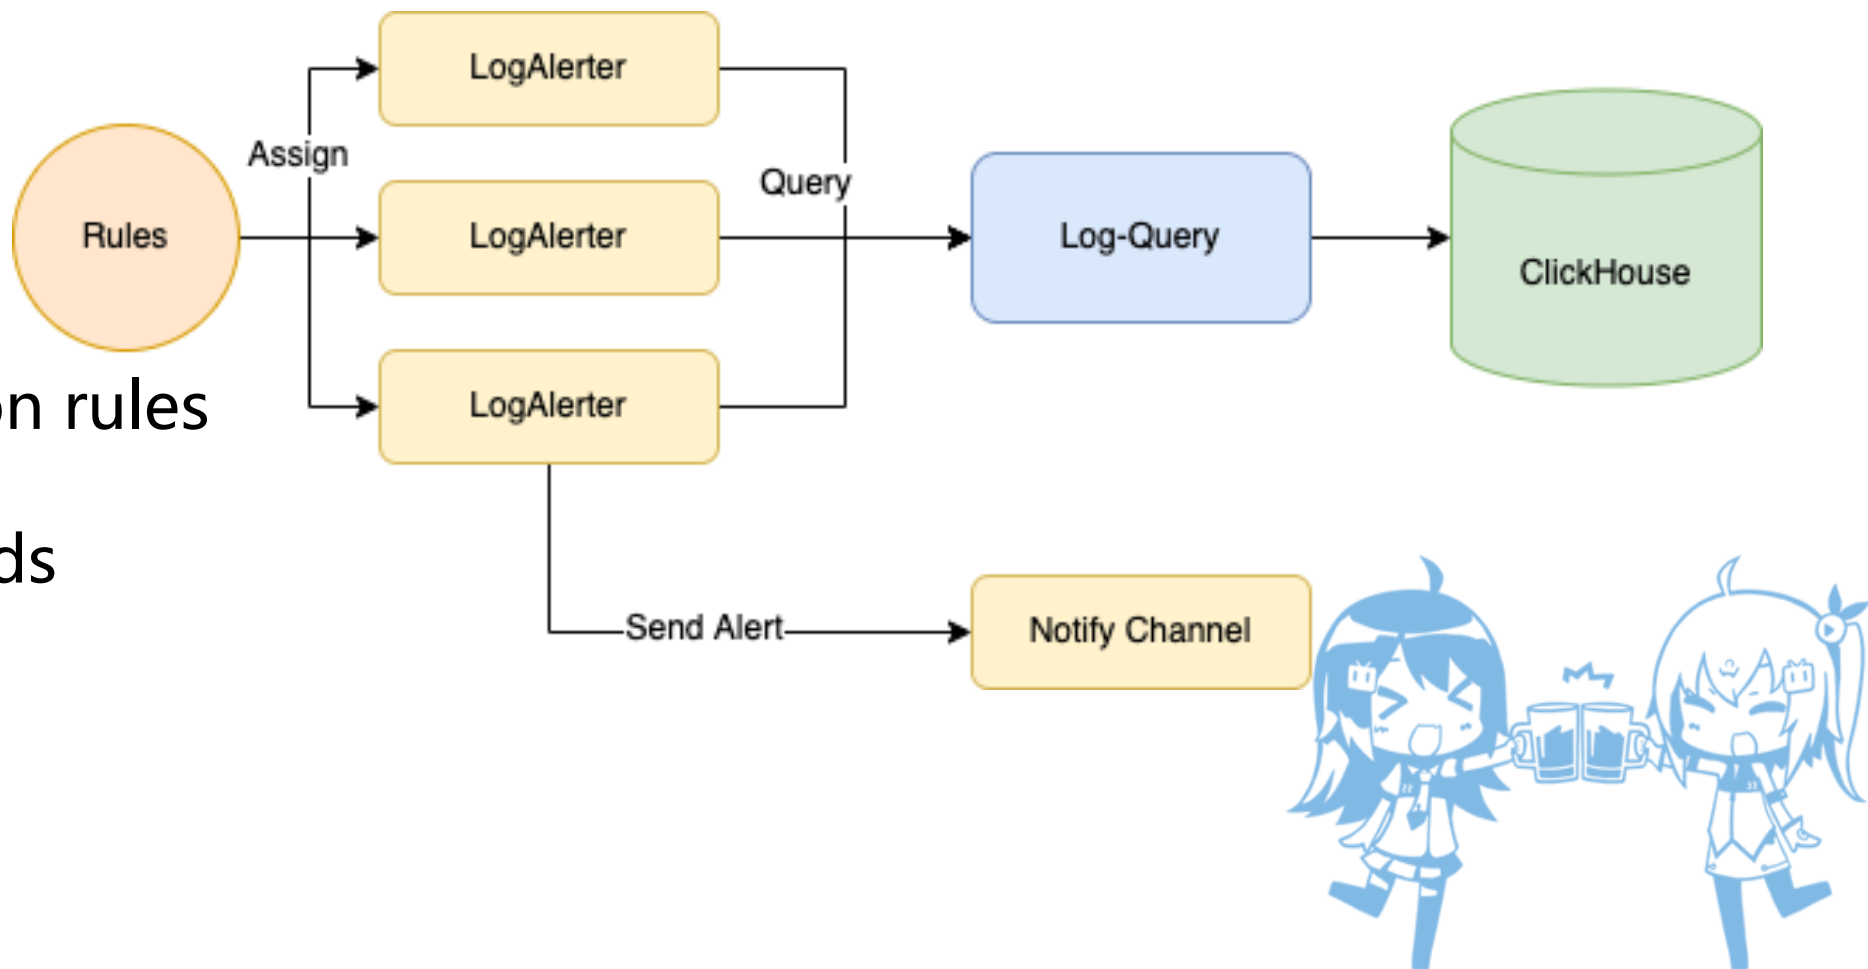
How we build



How we build

Log Alerter

- Distributed
- Various calculation rules
- Alert in 30 seconds





Optimization & Customization of ClickHouse for Log Analysis





Challenges for ClickHouse

Data writing:

- Large number of applications
- Low write latency
- Dynamic schema for private attribute
- Specify codec to save more disk

Data querying:

- Query based on application use different conditions
- Clickhouse native map type is underperforming



Overall Design

- Each application has a separate table

We can specify different settings based on application at table level.

1. Skipping indexes
2. Max implicit columns size

- Storage policy for hot and cold data

Specify storage policy in table DDL

1. NVME for hot data
2. SATA for cold data
3. ZSTD codec better than default codec (ZSTD can save up to 60% more space in some scenarios)

- Use Map type for dynamic attribute



Implicit Map

Native Map Type

1. Tuple of Arrays
2. Queries read redundant data
3. Skipping indexes don't support Map type

Array of Keys Array of values

k1	v1
k2	v2
k3	v3
k4	v4
k5	v5
k6	v6
k7	v7
...	...

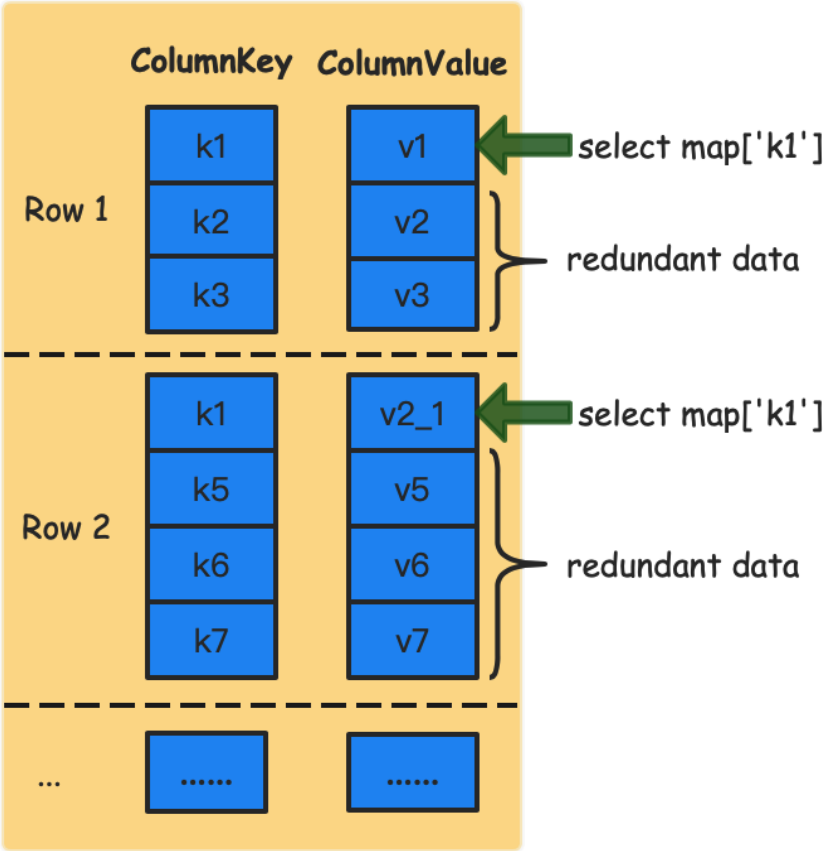


Implicit Map

Redundant data read

```
SELECT map['k1'] FROM table
```

- 2 arrays will be read
- All elements of map will be read
- Every rows of 2 arrays will be read





Skip Indices for Map

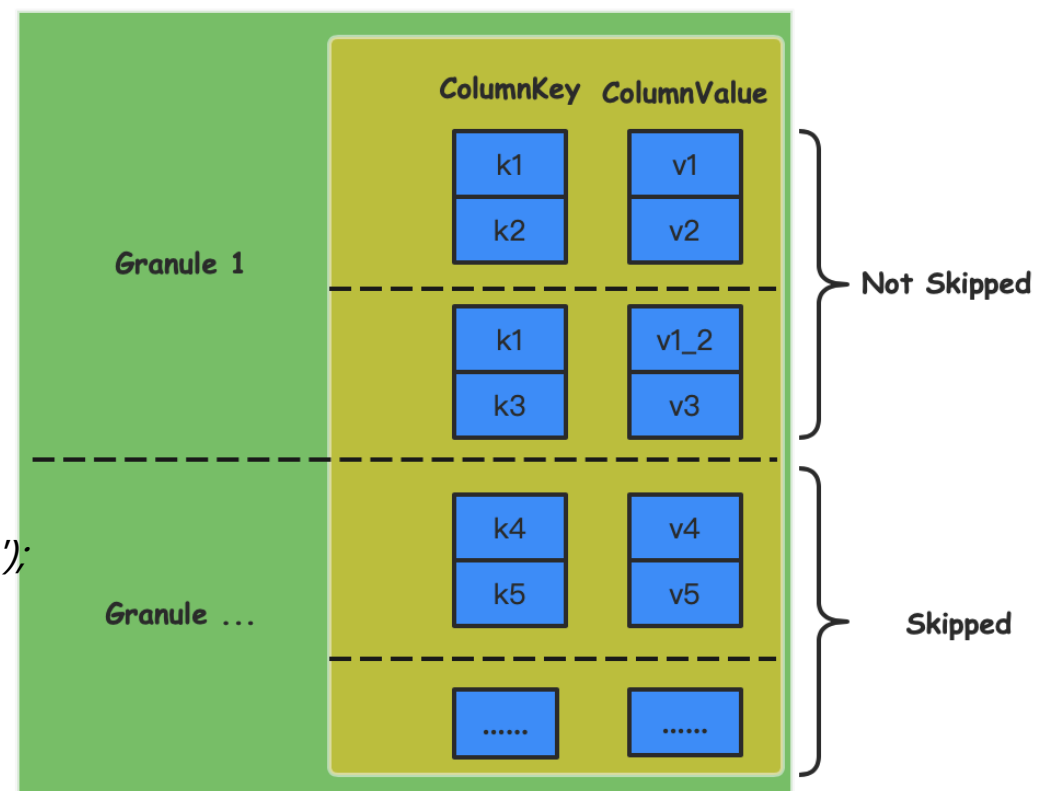
Skip Indices for Map

```
CREATE TABLE bloom_filter_map  
(  
  `id` UInt32,  
  `map` Map(String, String),  
  INDEX map_index map TYPE tokenbf_v1(128, 3, 0) GRANULARITY 1  
)  
ENGINE = MergeTree  
ORDER BY id  
SETTINGS index_granularity = 2  
  
SELECT map['k1'] FROM bloom_filter_map WHERE mapContains(map, 'k1');
```



Avoid reading unnecessary granules

Still read all elements in map

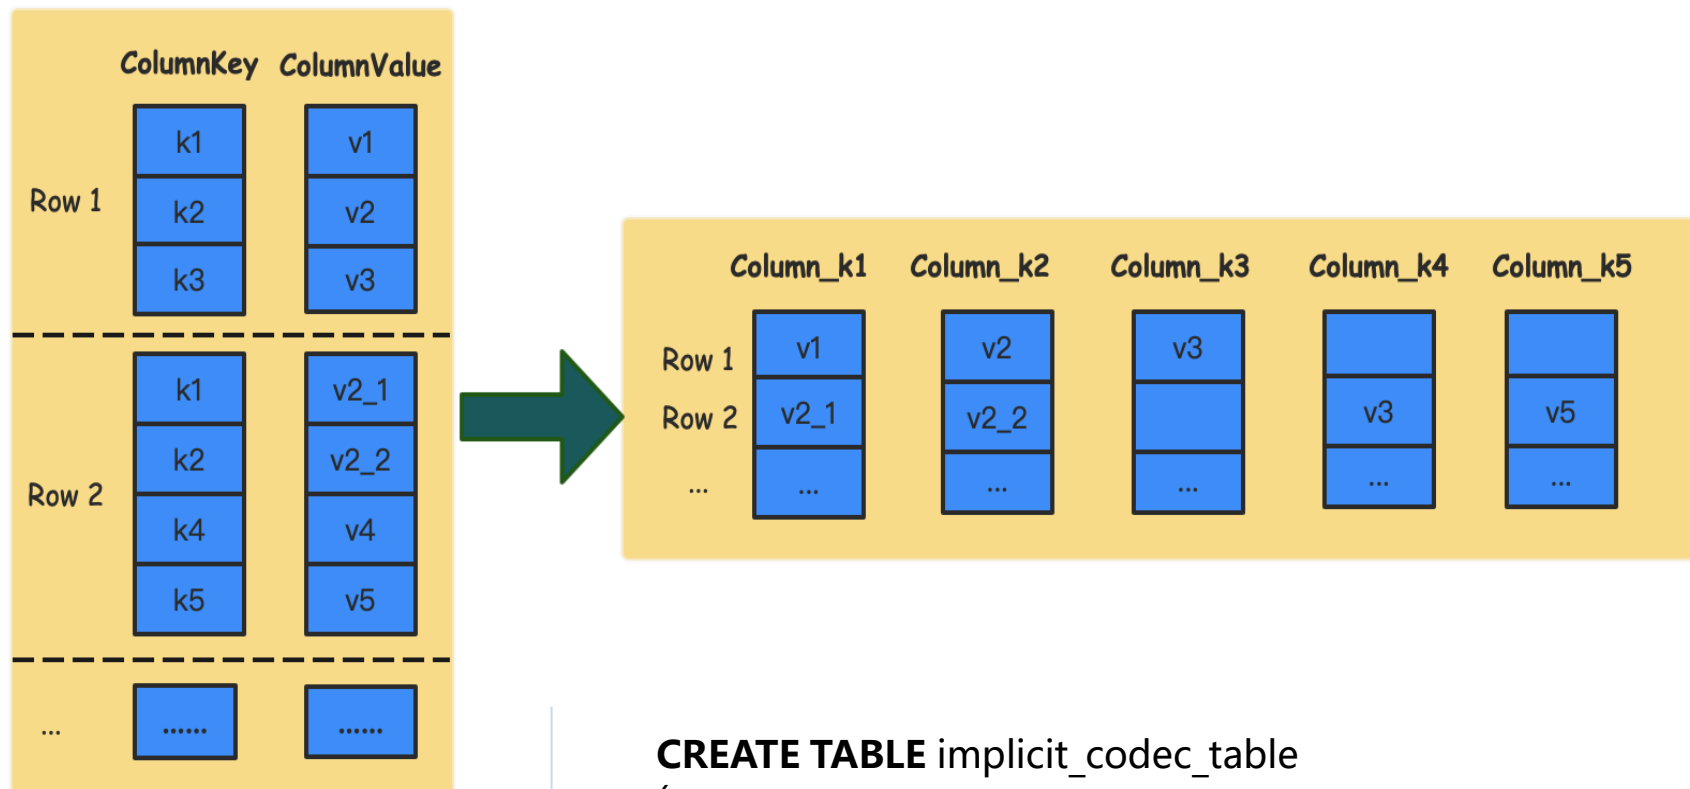




Implicit Map

What is Implicit Map?

Extract every unique key as a base column. This process is transparent to the users, so we called it as implicit columns.



```
CREATE TABLE implicit_codec_table
(
  `id` UInt32,
  `map` MapV2(String, Nullable(String)) CODEC(ZSTD(1)),
  `map2` MapV2(String, Nullable(String))
)
ENGINE = MergeTree
ORDER BY id
SETTINGS index_granularity = 8192
```



Implicit Map

Performance comparison

The data of two table is totally same

- Bytes on disk

Table Name	Rows	Bytes on disk
normal_map_table	10,000,000	3.19 GiB
implicit_map_table	10,000,000	3.11 GiB

- Query specified key

Table Name	Elapsed	Processed	RPS
normal_map_table	10.460 sec	5.07 GB	956.04 thousand rows/s 484.87 MiB/s
implicit_map_table	1.450 sec	551.34 MB	16.89 million rows/s 380.14 MiB/s.

Table DDL:

```
CREATE TABLE test.implicit_map_table
(
  `id` UInt64,
  `map` MapV2(String, String)
)
ENGINE = MergeTree
ORDER BY id
```

Query:

```
SELECT uniqExact(op)
FROM (
  SELECT map['OnePiece'] AS op
  FROM
    normal_map_table
)
```



Implicit Map

select all for Implicit Map

We have tried:

- Construct Map data from all implicit columns during query process
- We create implicit columns and also store native map


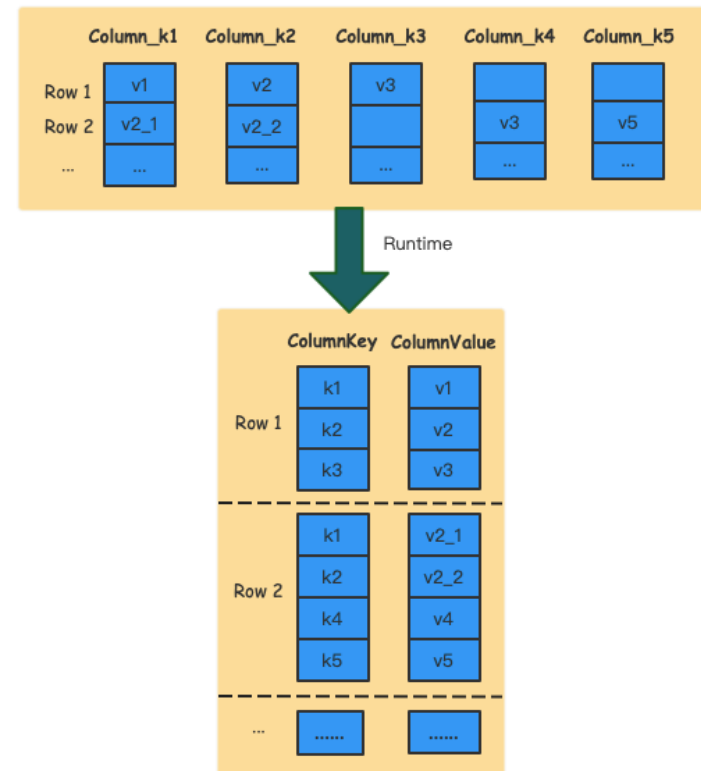
Query with specified key will read implicit column.
If users want get whole map, it will read native map column as usual.

- Specify every key

`SELECT map FROM map_table`

`SELECT map['k1'], map['k2'] ... map['kn'] FROM map_table`

We can get all implicit column name from system.parts.



Thanks

