

Metoda's journey into ClickHouse

Metoda, Munich, 2022-07-27

about me



Milovan Zogovic

head of engineering at metoda (Munich)

- ~20 years of experience in IT
- ruby developer by passion
- handling data at scale

business Facts & Numbers

We are the expert for insight-driven growth.



Founded: 2012



HQ: Munich



coverage:
global



High-quality, global
e-commerce **data**
& insights



Smart technologie for
Amazon (Ads) **process**
automation &
e-commerce **monitoring**



50+ industry experts
& certified Amazon
Ads team




Company values:
Drive Success Together



Awards:



Service
Provider
Network



tech Facts & Numbers

- processing around 10TB of data per day
- running 50 production stacks
- in-house data management solutions

Why Clickhouse?

A "typical" app

- Custom date ranges
- Customizable chart
- Filtering by any column
- Sorting by any column
- Aggregated facts as columns



- Unlimited history
- User changes take effect immediately
- Dashboard loads fast

A "typical" implementation

- Using existing OLTP database (postgres in our case)
- Joining facts and dimensions
- Aggregating, filtering and sorting done at database level
- ...

it works for a while, but things starts to get slow

A "typical" next steps

- Understanding why is it getting slower
- Profiling and fine tuning database queries
- Scaling up the database server (vertically and/or horizontally)
- Partitioning the data
- Optimizing the data ingestion
- ...

Are we still using the right technology?

What are some of the alternatives?

we did some research

- AWS Athena - way to slow for our use case
- AWS Redshift - faster, but very expensive to scale
- Elasticsearch - really fast, but no joins
- ClickHouse - seems to fulfil the requirements

Benchmark

Setup

- ~10GB of data, 150M rows facts, 1M rows dimensions
- Single node setup (with 2GB RAM)
- Using top accounts (by number of facts)
- A random 30 day interval in past 6 months
- 10 concurrent users
- Requests every 2-3 seconds
- 5 minutes benchmark time

Benchmark

```

checks.....: 100.00% ✓ 1121      ✗ 0
data_received.....: 958 kB  3.2 kB/s
data_sent.....: 1.4 MB  4.8 kB/s
http_req_blocked.....: avg=2.66ms  min=2.62µs  med=4.38µs  max=208.73ms  p(90)=9.16µs  p(95)=11.48µs
http_req_connecting.....: avg=801.46µs min=0s      med=0s      max=50.02ms  p(90)=0s      p(95)=0s
http_req_duration.....: avg=195.89ms min=107.73ms med=184.57ms max=795.31ms p(90)=253.33ms p(95)=285.7ms
  { expected_response:true }...: avg=195.89ms min=107.73ms med=184.57ms max=795.31ms p(90)=253.33ms p(95)=285.7ms
http_req_failed.....: 0.00% ✓ 0      ✗ 1121
http_req_receiving.....: avg=1.52ms  min=30.24µs med=972.84µs max=14.9ms  p(90)=3.58ms  p(95)=5.3ms
http_req_sending.....: avg=35.32µs min=14.09µs med=27.26µs  max=157.39µs p(90)=61.82µs p(95)=70.2µs
http_req_tls_handshaking.....: avg=978.6µs min=0s      med=0s      max=67.86ms p(90)=0s      p(95)=0s
http_req_waiting.....: avg=194.33ms min=103.78ms med=183.39ms max=795.23ms p(90)=251.83ms p(95)=285.62ms
http_reqs.....: 1121  3.711205/s
iteration_duration.....: avg=2.68s   min=2.12s   med=2.69s   max=3.66s   p(90)=3.08s   p(95)=3.14s
iterations.....: 1121  3.711205/s
vus.....: 2      min=2      max=10
vus_max.....: 10     min=10     max=10

```

Benchmark

Outcome

Results

- Average response times of 200ms on complex queries
- Unfair tests (limited CPU, limited RAM)

Considerations

- Benchmarks ran against fully sorted immutable data
- Data changes will affect performance (versioning, merging, sorting)

Next steps

Getting the data in

- Fact data loaded less frequently (in batches)
 - data resides on S3, and loaded directly from there into clickhouse
- Dimension data loaded more frequently
 - source application streams the changes
 - changes buffered and flushed every X seconds
- Versioning via ReplacingMergeTree
 - handles both "update" and "delete" scenario
- View on top
 - hides away complexity from the outside world

Other considerations

- Data management
 - data catalog (what data is there, who produces it, and how do we access it)
 - dependency management (who consumes what data)
- Making data available outside clickhouse
 - ETL processes
 - AWS integration
- Managing production environment (cluster setup, scaling, backups, availability.. etc)

Thank you!